# Model Match

An informed first model choice, grounded in the task you are building.

## 01 The user and problem

A full-stack developer is building their first AI feature on Fireworks: routing support tickets. They can integrate an API, but lack a test dataset and a method for deciding which model is good enough. A general benchmark does not establish how a model will handle their inbox. **Hypothesis:** guided, inspectable evidence helps them choose a starting model and understand its limitations.

## 02 The product experience

Review a prepared inbox and its expected routes; set minimum correct routes, a median response-time limit, a monthly budget and a priority; then compare three candidates. Inspect specific mistakes and leave with a provisional recommendation, a decision report, labeled test cases and optional starter code. A second inbox introduces mixed intent and missing context.

## 03 Key decisions and tradeoffs

**One task, three candidates:** keeps the first decision manageable. **Requirements before ranking:** filter on all three limits, then rank by cost, speed or accuracy. No qualifying model means no winner. **Reviewed labels:** developers can correct expected routes and see scores recalculate. **Visible failures:** ticket-level evidence explains the recommendation. **Simulation:** all outputs, timings and costs are invented and labeled; real model names are contextual. This proves the interaction, not model performance. No API key, upload flow, account system or universal score is required.

## 04 What I would validate next

Test with five developers new to model selection. Ask each to choose a model, explain a rejected alternative and identify what the result cannot establish. Observe whether they understand the requirements and simulation without coaching. Then replace fixtures with live inference and evaluate representative, human-labeled inbox data. Check repeatability, critical routing failures, token-based cost and latency variation; compare with browsing a model catalog.

## 05 How I would measure success

**Primary:** proportion of users who make an evidence-backed choice and correctly describe one limitation without help. **Usability:** completion rate and time from opening the tool to an explained choice. **Activation:** successful application API call after a real evaluation. **Guardrails:** rate of mistaking simulated results for real evidence, and critical-error rate on held-out tickets. Initial usability target: at least 4 of 5 participants complete and explain the decision unaided; this is directional, not statistical proof.